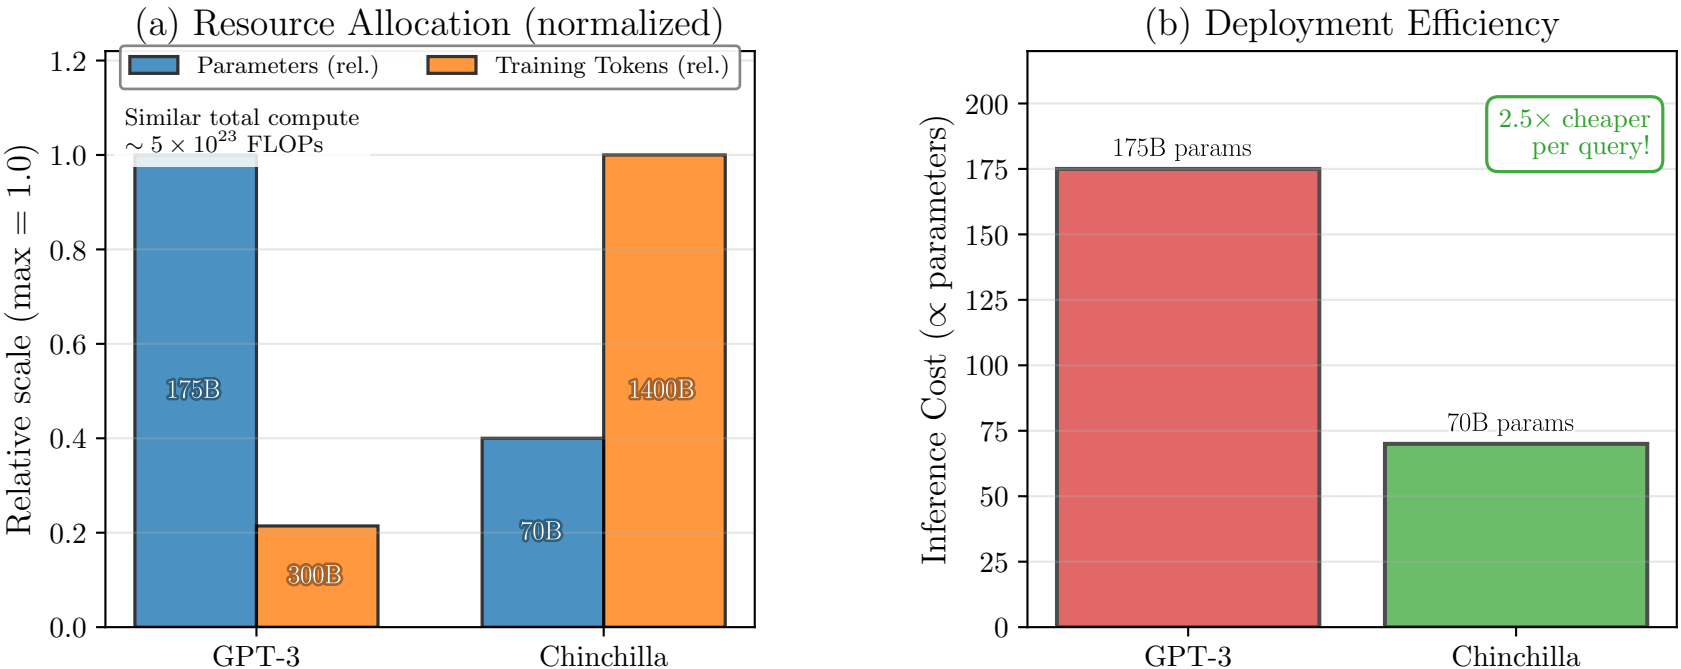
(a) Resource Allocation (normalized)
Parameters (rel.)
Training Tokens (rel.)
Similar total compute
$\sim 5 \times 10^{23}$ FLOPs
Relative scale (max = 1.0)
1.2
1.0
0.8
0.6
0.4
0.2
0.0
175B
300B
70B
1400B
GPT-3
Chinchilla
(b) Deployment Efficiency
Inference Cost ($\propto$ parameters)
200
175
150
125
100
75
50
25
0
175B params
70B params
2.5× cheaper
per query!
GPT-3
Chinchilla